

GPU rigid body simulation using OpenCL

Introduction

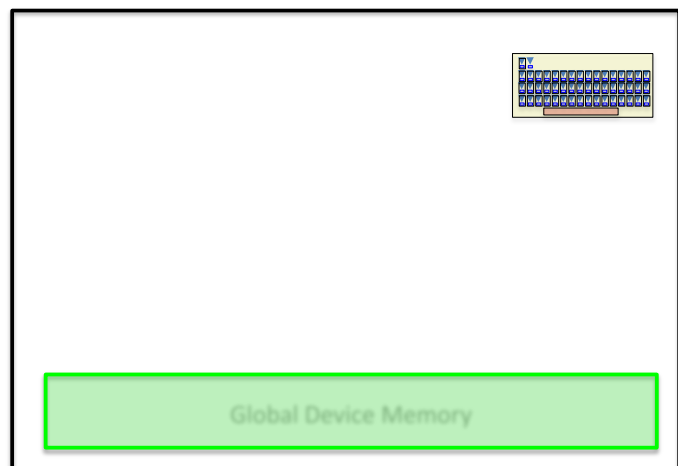
Bullet 2.x Refactoring

-
-
-
-
-

Bullet 3.x Full Rewrite

Getting started with OpenCL

OpenCL terminology



Our first OpenCL kernel

```
typedef struct
{
    float4 m_pos;
    float4 m_linVel;
} Body;

void integrateTransforms (Body* bodies, int nodeID, float timeStep)
{
    for (int nodeID = 0; nodeID < numBodies; nodeID++)
    {
        if (bodies[nodeID].m_invMass != 0.f)
        {
            bodies[nodeID].m_pos += bodies[nodeID].m_linVel * timeStep;
        }
    }
}
```

```
__kernel void integrateTransformsKernel( __global Body* bodies, const int numNodes, float
timeStep)
{
    int nodeID = get_global_id(0);
    if ( nodeID < numNodes && (bodies[nodeID].m_invMass != 0.f))
    {
        bodies[nodeID].m_pos += bodies[nodeID].m_linVel * timeStep;
    }
}
```

Porting existing code to OpenCL

Replace C++ by C


Move data to contiguous memory

Replace pointers by indices

```
struct btTransform
{
    btMatrix3x3    m_basis;
    btVector3      m_position;
};
class btRigidBody : public btCollisionObject
{
    btMatrix3x3    m_inverseInertiaWorld;
    btVector3      m_linearVelocity;
    btVector3      m_angularVelocity;
    btScalar       m_mass;
    . . .
};
class btCollisionObject
{
    btTransform     m_worldTransform;
    btCollisionShape* m_collisionShape;
    . . .
};
```

```
struct b3RigidBody
{
    b3Vector3      m_position;
    b3Quaternion    m_orientation;
    int            m_collidableIndex;
    . . .
};
struct b3Collidable
{
    int m_shapeType;
    int m_shapeIndex;
};
```

Exploiting GPU hardware



Dealing with branchy code/thread divergence

```
kernel void branchyKernel (. . .)
{
    if (conditionA)
    {
        someCodeA(. . .);
    } else
    {
        someCodeNotA(. . .);
    }
}
```

Use Parallel Primitives

-
-
-

Use Local Memory

-
-
-

Barrier synchronization

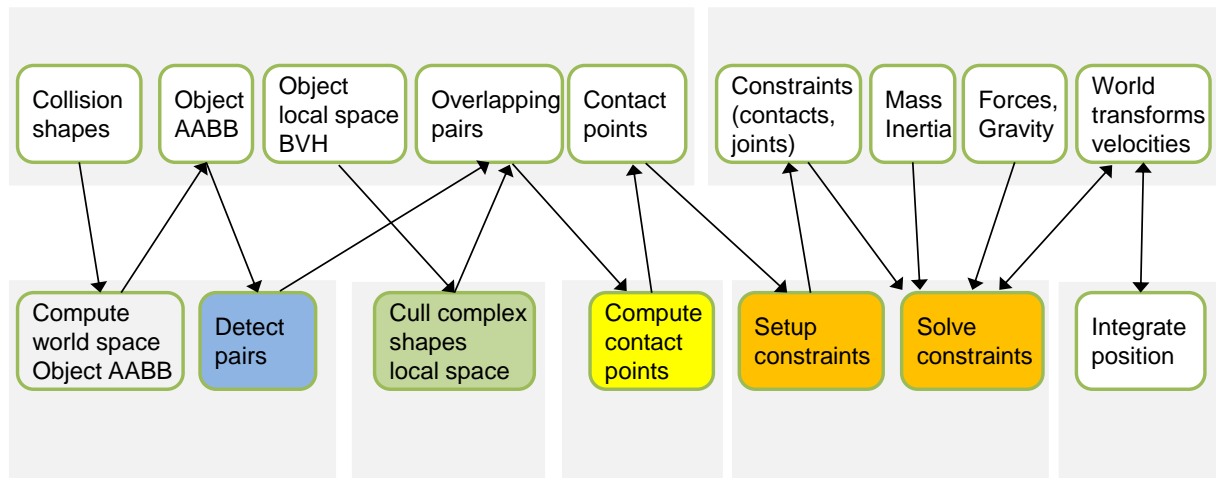
-
-
-

Atomics


GPU rigid body simulation

Rigid body introduction

The rigid body pipeline



Computing the object AABBs



GPU overlapping pair detection

```

void computePairsKernelBruteForce (const btAabbCL* aabbs, volatile __global int2*
pairsOut,volatile __global int* pairCount, int numObjects, int maxPairs)
{
    for (int i=0;i<numObjects;i++)
    {
        for (int j=i+1;j<numObjects;j++)
        {
            if (TestAabbAgainstAabb2GlobalGlobal(&aabbs[i],&aabbs[j]))
            {
                int2 myPair;
                myPair.x = aabbs[i].m_objectIndex;
                myPair.y = aabbs[j].m_objectIndex;
                int curPair = *pairCount;
                if (curPair<maxPairs)
                {
                    pairsOut[curPair] = myPair; //flush to main memory
                    pairCount++;
                }
            }
        }
    }
}


```


```

__kernel void computePairsKernelBruteForceKernel ( __global const btAabbCL* aabbs, volatile
__global int2* pairsOut,volatile __global int* pairCount, int numObjects, int maxPairs)
{
    int i = get_global_id(0);
    if (i>=numObjects)
        return;
    for (int j=i+1;j<numObjects;j++)
    {
        if (TestAabbAgainstAabb2GlobalGlobal(&aabbs[i],&aabbs[j]))
        {
            int2 myPair;
            myPair.x = aabbs[i].m_objectIndex;
            myPair.y = aabbs[j].m_objectIndex;
            int curPair = atomic_inc (pairCount);
            if (curPair<maxPairs)
            {
                pairsOut[curPair] = myPair; //flush to main memory
            }
        }
    }
}


```

•
•
•
•
•







GPU local space BVH culling for complex shapes




GPU contact computation




```
void computeContactsSAT(convexHullA,convexHullB,transformA,transformB)
{
    b3Vector3 satAxis;
    if (findSeparatingAxis(convexHullA,convexHullB,transformA,transformB))
    {
        if (clipHullHull(satAxis, convexHullA,convexHullB,transformA,transformB))
        {
            findClippingFacesKernel(. . .)
            clipFacesAndContact(. . .)
            contactReduction(...)
        }
    }
}
```



GPU parallel contact solving



```
void batchConstraints(constraints, int numConstraints)
{
    int batchIdx=0;
    while( numValidConstraints < numConstraints)
    {
        clear(bodiesUsed);
        for(int i=numValidConstraints; i<numConstraints; i++)
        {
            int bodyA = constraints[i].m_bodyA;
            int bodyB = constraints [i].m_bodyB;
            if (isAvailable(bodyA,bodyB,bodiesUsed))
            {
                markUsed(bodyA,bodyB,bodiesUsed);
                constraint[i].m_batchId = batchIdx;//assign the batch index
            }
        }
        batchIdx ++;
    }
}
```

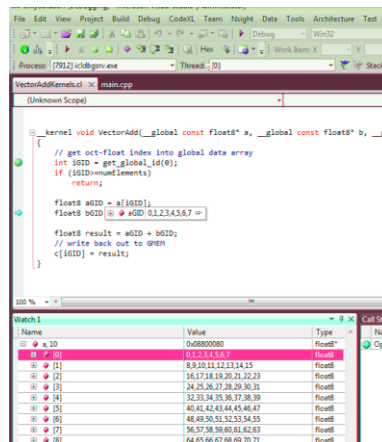


GPU parallel joint solving

Debugging and Performance Profiling

Debug on the CPU

Intel OpenCL debugger



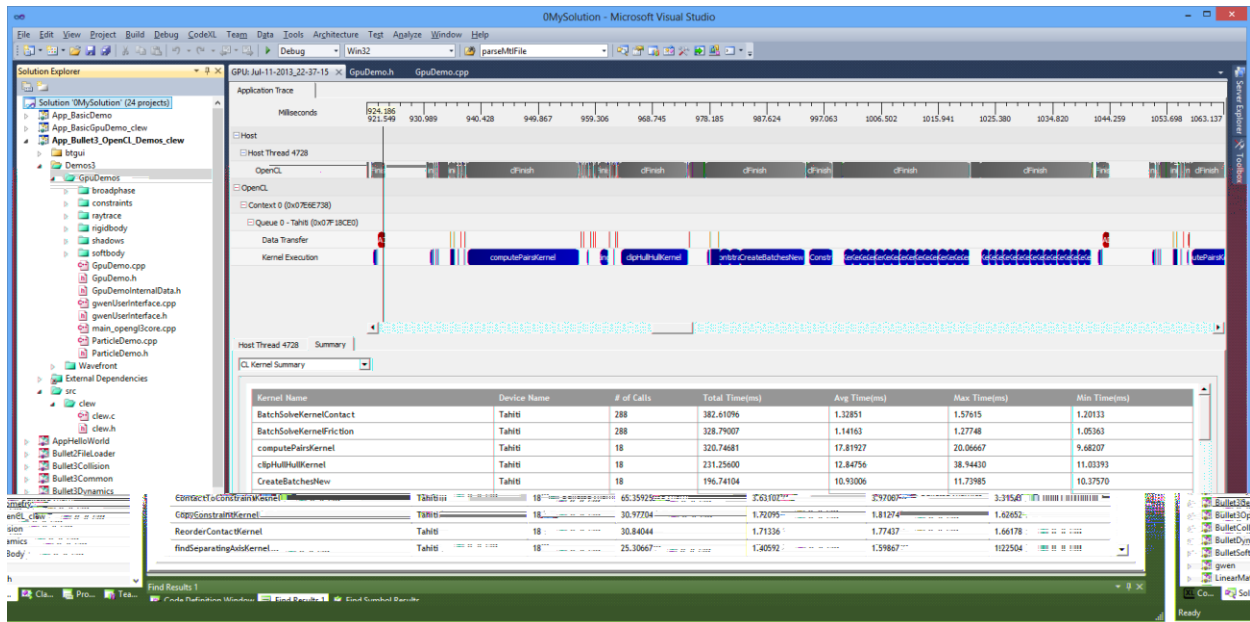
printf debugging

Debug buffers

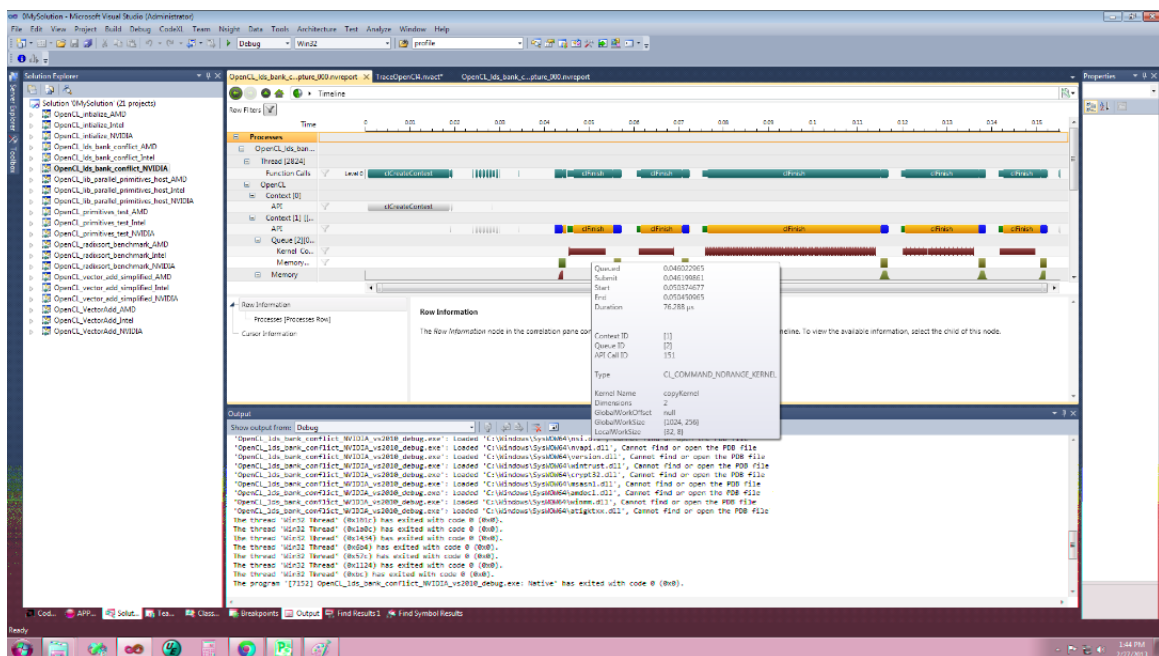
Debugging a frozen system

Profile Zones

CodeXL Performance Profiler



NVIDIA NSIGHT Profiler



OpenCL Tips and Tricks

Create your own OpenCL wrapper


Dynamically load OpenCL

Cache the precompiled OpenCL kernel binaries

Keep a host implementation of your kernel

Unit test an OpenCL kernel

References



Appendix A: Bullet 3.x Source code



Requirements


Building on Windows using Visual Studio

Building on Linux (or Mac OSX) using gcc

```
cd build3
./premake_linux64 gmake
cd gmake
make
```

Building on Max OSX using XCode

Usage



Benchmark

Feedback
